

Chapter 3

Naïve Bayes Learning

Contents

3 Naïve Bayes Learning	1
3.1 Direct Learning	2
3.2 Probabilistic Model	3
3.3 Probability Recap	3
3.4 Joint Distribution	3
3.5 Probability Distribution	4
3.5.1 Bernoulli Distribution	4
3.5.2 Binomial Distribution	4
3.5.3 Categorical Distribution	4
3.5.4 Multinomial Distribution	5
3.6 Independence	5
3.7 Bayes' Rule	5
3.8 Bayesian Learning	6
3.9 Maximum APosteriori Estimate	6
3.10 Maximum Likelihood Estimate	7
3.11 Bayesian Classifier	7
3.12 Naïve Bayes Classifier	8
3.13 Estimating Probabilities	8
3.14 Laplace Smoothing	8
3.15 Continuous Features	9
3.16 Gaussian Naïve Bayes	9
3.17 Bayesian Belief Network	9
3.18 Training Bayesian Classifier	10
3.19 Text Classification	10
3.20 Evaluating Classifiers	11
3.21 Precision, Recall, F-Measure	11
3.22 ROC Curve	12
3.23 AUC	12
3.24 Naïve Bayes: Two Classes	13

3.1 Direct Learning

- Consider a **distribution** D

- X - Instance space, Y - Set of labels. (e.g. ± 1)
- Given a sample $\{(\mathbf{x}, y)\}_1^n$ and a loss function $L(\mathbf{x}, y)$, find a hypothesis $h \in H$ that minimizes $\sum_{i=1 \dots n} L(h(\mathbf{x}_i), y_i)$.

Table 3.1: Losses

0 - 1 loss:	$L(h(\mathbf{x}), y) = 1, h(\mathbf{x}) \neq y$ otherwise $L(h(\mathbf{x}), y) = 0$
L_2 Loss:	$L(h(\mathbf{x}), y) = (h(\mathbf{x}) - y)^2$
Hinge Loss:	$L(h(\mathbf{x}), y) = \max\{0, 1 - yh(\mathbf{x})\}$
Exponential Loss:	$L(h(\mathbf{x}), y) = e^{-yh(\mathbf{x})}$

3.2 Probabilistic Model

Paradigm:

- Learn a probability distribution of the dataset.
- Use it to estimate which outcome is more likely.

Instead of learning $h : X \rightarrow Y$, learn $P(Y|X)$.

- Estimate probability from data
 - Maximum Likelihood Estimate (MLE)
 - Maximum A posteriori Estimation (MAP)

3.3 Probability Recap

$$\begin{aligned}
 0 &\leq P(A) \leq 1 \\
 P(\text{true}) &= 1, P(\text{false}) = 0 \\
 P(A \vee B) &= P(A) + P(B) - P(A \wedge B) \\
 P(A|B) &= \frac{P(A \wedge B)}{P(B)}
 \end{aligned}$$

3.4 Joint Distribution

Making a joint distribution of d variables

- Make a truth table listing all combinations of values of your variables (if there are d boolean variables then the table will have 2^d rows)

- For each combination of values, say how probable it is.
- The probability must sum up to 1.

Once we have the Joint Distribution, we find probability of any logical expression involving these variables.

$$\begin{aligned}
 P(E) &= \sum_{\text{rows matching } E} P(\text{row}) \\
 P(E_1 \mid E_2) &= \frac{E_1 \wedge E_2}{P(E_2)} \\
 &= \frac{\sum_{\text{rows matching } E_1 \text{ and } E_2} P(\text{row})}{\sum_{\text{rows matching } E_2} P(\text{row})}
 \end{aligned} \tag{3.1}$$

3.5 Probability Distribution

3.5.1 Bernoulli Distribution

Random Variable X takes values $\{0, 1\}$ such that

$$P(X = 1) = p = 1 - P(X = 0)$$

Example: Tossing a coin.

3.5.2 Binomial Distribution

Random Variable X takes values $\{1, 2, \dots, n\}$ representing the number of successes $X = 1$ in n Bernoulli trials.

$$P(X = k) = f(n, p, k) = C_n^k p^k (1 - p)^{n-k}$$

Example: Tossing a coin n times.

3.5.3 Categorical Distribution

Random Variable X takes on values in $\{1, 2, \dots, k\}$ such that

$$P(X = i) = p_i \text{ and } \sum_1^k p_i = 1$$

Example: Rolling a die.

3.5.4 Multinomial Distribution

- Let the random variables $X_i (i = 1, 2, \dots, k)$ indicates the number of times outcome i was observed over the n trials.
- The vector $X = (X_1, X_2, \dots, X_k)$ follows a multinomial distribution (n, p) where $p = (p_1, p_2, \dots, p_k)$ and $\sum_1^k = 1$

$$f(x_1, x_2, \dots, x_k, n, p) = P(X_1 = x_1, X_2 = x_2, \dots, X_k = x_k)$$

$$= \frac{n!}{x_1! \times x_2! \times \dots \times x_k!} \times p_1^{x_1} \times p_2^{x_2} \times \dots \times p_k^{x_k} \text{ where } \sum_{i=1}^k x_i = n$$

Example: Rolling a die n times.

3.6 Independence

When two events do not affect each other's probabilities, they are called independent events

$$A \perp\!\!\!\perp B \Leftrightarrow P(A \wedge B) = P(A) \times P(B)$$

$$\Leftrightarrow P(A | B) = P(A)$$

The conditional independence of events A and B , given C is:

$$A \perp\!\!\!\perp B | C \Leftrightarrow P(A | B, C) = \frac{P(A \wedge B | C)}{P(B | C)} = \frac{P(A | C) \times P(B | C)}{P(B | C)}$$

$$= P(A | C)$$

3.7 Bayes' Rule

$$P(A|B) = \frac{P(B|A) \times P(A)}{P(B)} \quad (3.2)$$

where A and B are events and $P(B) \neq 0$. Applying Bayes' rule for machine learning –

$$P(\text{hypothesis} | \text{evidence}) = \frac{P(\text{evidence} | \text{hypothesis}) \times P(\text{hypothesis})}{P(\text{evidence})} \quad (3.3)$$

3.8 Bayesian Learning

- Goal: find the **best hypothesis** from some space H of **hypotheses**, given the observed data (**evidence**) D .
- Define the **most probable hypothesis** in H to be the **best**.
- In order to do that, we need to **assume** a **probability distribution** over the **class** H .
- In addition, we need to know something about the **relation** between the **evidence** and the **hypotheses**.

$P(h)$ – **Prior Probability** of the **hypothesis** h . Reflects the background knowledge, before data is observed.

$P(D)$ – **Probability** that this sample of the **data** is **observed**.

$P(D|h)$ – Probability of **observing** the **sample** D , given that **hypothesis** h is the **target**, also referred to as **likelihood**.

$P(h|D)$ – **Posterior probability** of h . The **probability** that h is the **target**, given that D has been **observed**.

- $P(h|D)$ **increases** with $P(h)$ and $P(D|h)$.
- $P(h|D)$ **decreases** with $P(D)$.

3.9 Maximum APosteriori Estimate

$$P(h|D) = \frac{P(D|h) \times P(h)}{P(D)} \quad (3.4)$$

- The **learner** considers a **set of candidate hypotheses** H (models) and attempts to find the **most probable** one $h \in H$, given the observed data.
- Such maximally probable hypothesis is called **maximum a posterior estimate** (MAP). Bayes theorem is used to compute it:

$$\begin{aligned} h_{MAP} &= \arg \max_{h \in H} P(h|D) \\ &= \arg \max_{h \in H} \frac{P(D|h) \times P(h)}{P(D)} \\ &= \arg \max_{h \in H} P(D|h) \times P(h) \end{aligned}$$

3.10 Maximum Likelihood Estimate

- We may assume that *a priori*, hypotheses are equally probable.

$$P(h_i) = P(h_j) \forall h_i, h_j \in H$$

- With that assumption, we can treat $\frac{P(h)}{P(D)}$ as a constant. We get the *maximum likelihood estimate* (MLE):

$$\begin{aligned} h_{MLE} &= \arg \max_{h \in H} \frac{P(D|h) \times P(h)}{P(D)} \\ &= \arg \max_{h \in H} P(D|h) \end{aligned}$$

- Here we just *look for* the hypothesis that *best explains* the *data*.

3.11 Bayesian Classifier

- $f: \mathbf{X} \rightarrow Y$ where, instances $\mathbf{x} \in \mathbf{X}$ is a collection of inputs –

$$\mathbf{x} = (x_1, x_2, \dots, x_n)$$

- Given an example, *assign* it the *most probable* value in Y .

$$\begin{aligned} y_{MAP} &= \arg \max_{y_j \in Y} P(y_j | x) \\ &= \arg \max_{y_j \in Y} P(y_j | x_1, x_2, \dots, x_n) \\ &= \arg \max_{y_j \in Y} \frac{P(x_1, x_2, \dots, x_n | y_j) P(y_j)}{P(x_1, x_2, \dots, x_n)} \\ &= \arg \max_{y_j \in Y} P(x_1, x_2, \dots, x_n | y_j) P(y_j) \end{aligned} \tag{3.5}$$

- Given the training data, we have to *estimate* the two terms.
- Estimating $P(y)$ is easy, e.g., under the *binomial distribution assumption*, *count* the number of *times* y appears in the training data.
- However, it is *not feasible* to estimate $P(x_1, x_2, \dots, x_n | y)$
- In this case, we have to *estimate* for each *target* value, the *probability* of *each instance* (some of which might now ever occur).
- In order to use a Bayesian classifiers in practice, we need to *make assumptions* that will *allow* us to *estimate* these quantities.

3.12 Naïve Bayes Classifier

Assumption: Input feature values are independent, given the target value.

$$\begin{aligned}
 P(x_1, x_2, \dots, x_n | y_j) &= P(x_1 | y_j) \times P(x_2, \dots, x_n | y_j) \\
 &= P(x_1 | y_j) \times P(x_2 | y_j) \times P(x_3, \dots, x_n | y_j) \\
 &= P(x_1 | y_j) \times P(x_2 | y_j) \times P(x_3 | y_j) \times \dots \times P(x_n | y_j) \\
 &= \prod_{i=1}^n P(x_i | y_j)
 \end{aligned} \tag{3.6}$$

$$\begin{aligned}
 Y_{NB} &= \arg \max_{y_j \in Y} P(x_1, x_2, \dots, x_n | y_j) P(y_j) \\
 &= \arg \max_{y_j \in Y} P(y_j) \prod_{i=1}^n P(x_i | y_j)
 \end{aligned} \tag{3.7}$$

3.13 Estimating Probabilities

How do we estimate $P(x_i | y)$?

$$P(x_i | y) = \frac{\text{number of } x_i \text{ labeled as } y}{\text{total number of label } y} = \frac{n_i}{n} \tag{3.8}$$

Sparsity of data is a problem –

- If n is **small**, the estimate is not accurate.
- If $n_i = 0$, we will never accurately predict Y if an **instance** that **never appeared** in the **training** appears in the test data.

3.14 Laplace Smoothing

$$P(x_i | y) = \frac{n_i + \alpha}{n + \alpha d} \tag{3.9}$$

- Also known as **additive smoothing**.
- $\alpha > 0$ is a smoothing parameter.
- d is the **dimension** of the input.

3.15 Continuous Features

- Assume $P(x_i|y)$ has a Gaussian (normal) distribution.
- It is a continuous distribution with probability density function:

$$p(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \quad (3.10)$$

- μ is the mean of the distribution.
- σ^2 is the variance of the distribution.
- x is a continuous variance ($-\infty \leq x \leq \infty$)

3.16 Gaussian Naïve Bayes

Table 3.2: Naïve Bayes Example

X_1	X_2	X_3	Y
2	3	1	1
-1.2	2	0.4	1
1.2	0.3	0	0
2.2	1.1	0	1

Compute the mean and standard deviation to estimate the likelihood.

$$\begin{aligned} \mu_1 &= E[X_1 | Y = 1] = \frac{2 + (-1.2) + 2.2}{3} = 1 \\ \sigma_1^2 &= E[(X_1 - \mu_1)^2 | Y = 1] = \frac{(2 - 1)^2 + (-1.2 - 1)^2 + (2.2 - 1)^2}{3} = 2.43 \\ P(x_1 | Y = 1) &= \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x_1 - \mu_1)^2}{2\sigma^2}} = \frac{1}{3.91} e^{-\frac{(x_1 - 1)^2}{4.86}} \end{aligned}$$

3.17 Bayesian Belief Network

- Naïve Bayes classifier works with the assumption that the values of the input features are conditionally independent given the target value.
- This assumption dramatically reduces the complexity of learning the target function.
- Bayesian Belief Network describes the probability distribution governing a set of variables by specifying a set of conditional independence assumptions along with a set of conditional probabilities. Conditional independence assumptions here apply to subsets of the variables.

$$P(x_1, x_2, \dots, x_l | x_1', x_2', \dots, x_m', y_1, y_2, \dots, y_n) = P(x_1, x_2, \dots, x_l | y_1, y_2, \dots, y_n)$$

3.18 Training Bayesian Classifier

During [training](#), typically [log-space](#) is used.

$$\begin{aligned}
 y_{NB} &= \arg \max_y \left[\log P(y) \prod_{i=1}^n P(x_i|y) \right] \\
 &= \arg \max_y \left[\log P(y) + \sum_{i=1}^n \log P(x_i|y) \right]
 \end{aligned}$$

3.19 Text Classification

Algorithm 3.1 Text-based Naïve Bayes Classification

```

1: function TRAIN-NAIVE-BAYES( $D, C$ ) returns  $\log P(c)$  and  $\log P(w|c)$ 
2:   for all class  $c \in C$  do ▷ Calculate  $P(c)$  terms
3:      $N_{doc} \leftarrow$  number of documents in  $D$ 
4:      $N_c \leftarrow$  number of documents from  $D$  in class  $c$ 
5:      $\logprior[c] \leftarrow \log \frac{N_c}{N_{doc}}$ 
6:      $V \leftarrow$  vocabulary of  $D$ 
7:      $bigdoc[c] \leftarrow$  APPEND( $d$ ) for  $d \in D$  with class  $c$ 
8:     for all word  $w$  in  $V$  do ▷ Calculate  $P(w|c)$  terms
9:        $COUNT(w, c) \leftarrow$  # of occurrences of  $w$  in  $bigdoc[c]$ 
10:       $\loglikelihood[w, c] \leftarrow \log \frac{COUNT(w, c) + 1}{\sum_{w' \text{ in } V} (COUNT(w', c) + 1)}$ 
11:    end for
12:  end for
13:  return  $\logprior, \loglikelihood, V$ 
14: end function

```

Algorithm 3.2 Test Naïve Bayes

```

1: function TEST-NAIVE-BAYES( $testdoc, \logprior, \loglikelihood, C, V$ ) returns best  $c$ 
2:   for all class  $c \in C$  do
3:      $sum[c] \leftarrow \logprior[c]$ 
4:     for all position  $i$  in  $testdoc$  do
5:        $word \leftarrow testdoc[i]$ 
6:       if  $word \in V$  then
7:          $sum[c] \leftarrow sum[c] + \loglikelihood[word, c]$ 
8:       end if
9:     end for
10:  end for
11:  return  $\arg \max_c, sum[c]$ 
12: end function

```

The word **with** doesn't occur in the training set, so we drop it completely (we don't use unknown word models for Naïve Bayes)

3.20 Evaluating Classifiers

- **Gold Label** is the **correct** output **class** label of an input.
- **Confusion Matrix** is a table for **visualizing** how a **classifier performs** with respect to the gold labels, using two dimensions (system output and gold labels), and each cell labeling a set of possible outcomes.
- **True Positives** and **True Negatives** are **correctly classified** outputs belonging to the positive and negative class, respectively.
- **False Positives** and **False Negatives** are **incorrectly classified** outputs.

		<i>gold standard labels</i>		
		gold positive	gold negative	
<i>system output labels</i>	system positive	true positive	false positive	precision = $\frac{tp}{tp+fp}$
	system negative	false negative	true negative	
		recall = $\frac{tp}{tp+fn}$		accuracy = $\frac{tp+tn}{tp+fp+tn+fn}$

Figure 3.1: An example of a Confusion Matrix.

3.21 Precision, Recall, F-Measure

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}} \quad (3.11)$$

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}} \quad (3.12)$$

$$\text{F - measure} = F_{\beta} = \frac{(\beta^2 + 1)PR}{\beta^2 P + R} \quad (3.13)$$

3.22 ROC Curve

- A **receiver operating characteristic** curve (ROC curve) is a graphical plot that illustrates the **performance** of a **binary classifier model**.
- The **ROC curve** is the plot of the **true positive rate (recall)** (TPR) (3.12) against the **false positive rate** (FPR).

$$\text{FPR} = \frac{\text{false positives}}{\text{false positives} + \text{true negatives}} \quad (3.14)$$

- **ROC curve** plots **TPR vs. FPR** at different **classification thresholds**.
- **Classification threshold** is used to convert the **output** of a **probabilistic classifier** into class **labels**.
- The **threshold** determines the **minimum probability** required for a **positive class**.
- Lowering the classification threshold classifies more items as **positive**, thus increasing both False Positives and True Positives.




Figure 3.2: ROC Curve




Figure 3.3: ROC Curve with defined thresholds

3.23 AUC

- The Area Under the Curve (AUC) provides an **aggregate** measure of **performance** across all possible classification thresholds.
- Between two ROC curves plotted based on two learning models, the model with the higher AUC **learned better** than the other.

3.24 Naïve Bayes: Two Classes

- Naïve Bayes classifier gives a method for **predicting** the **most likely class** rather than an explicit class.
- In the case of two classes, $y \in \{0, 1\}$ we predict that $y = 1$ iff

$$\frac{P(y_j = 1) \times \prod_{i=1}^n P(x_i | y_j = 1)}{P(y_j = 0) \times \prod_{i=1}^n P(x_i | y_j = 0)} > 1$$

- $p_i = P(x_i | y_j = 1)$, $q_i = P(x_i | y_j = 0)$. Assuming Bernoulli Naïve Bayes,

$$\begin{aligned} & \frac{P(y_j = 1) \times \prod_{i=1}^n p_i^{x_i} (1 - p_i)^{1-x_i}}{P(y_j = 0) \times \prod_{i=1}^n q_i^{x_i} (1 - q_i)^{1-x_i}} > 1 \\ \Rightarrow & \frac{P(y_j = 1) \times \prod_{i=1}^n (1 - p_i)(p_i / (1 - p_i))^{x_i}}{P(y_j = 0) \times \prod_{i=1}^n (1 - q_i)(q_i / (1 - q_i))^{x_i}} > 1 \end{aligned}$$

Take logarithm; we predict $y = 1$ iff

$$\log \frac{P(y_j = 1)}{P(y_j = 0)} + \sum_i \log \frac{1 - p_i}{1 - q_i} + \sum_i \left(\log \frac{p_i}{1 - p_i} - \log \frac{q_i}{1 - q_i} \right) x_i > 0$$

- We get that Naïve Bayes is a **linear separator** with –

$$w_i = \log \frac{p_i}{1 - p_i} - \log \frac{q_i}{1 - q_i} = \log \frac{p_i(1 - q_i)}{q_i(1 - p_i)} \quad (3.15)$$

- NB classifier corresponds to a **linear classifier** if the likelihood is from **exponential family distributions** i.e., Bernoulli, binomial, Gaussian etc.
- In the case of two classes, we can say:

$$\log \frac{P(y_j = 1 | x)}{P(y_j = 0 | x)} = \sum_i \mathbf{w}_i \mathbf{x}_i + b$$

- but since $P(y_j = 1 | x) = 1 - P(y_j = 0 | x)$, we get:

$$P(y_j = 1 | x) = \frac{1}{1 + e^{-(\sum_i w_i x_i + b)}}$$

- This is simply the **logistic function**.