

Chapter 4

Logistic Regression

Contents

| | | |
|----------|---|----------|
| 4 | Logistic Regression | 1 |
| 4.1 | Generative Classifiers | 2 |
| 4.2 | Discriminative Classifiers | 2 |
| 4.3 | Generative vs Discriminative Classifiers | 3 |
| 4.4 | Learning a Logistic Regression Classifier | 3 |
| 4.5 | Logistic Regression | 4 |
| 4.6 | LR Example | 4 |
| 4.7 | Sentiment Classification | 5 |
| 4.8 | Training Logistic Regression | 6 |
| 4.9 | Cross-Entropy Loss | 6 |
| 4.10 | Minimizing Cross-Entropy Loss | 7 |
| 4.11 | Optimizing a Convex/Concave Function | 8 |
| 4.12 | Gradients | 8 |
| 4.13 | Gradient Descent for Logistic Regression | 8 |
| 4.14 | Learning Rate | 9 |
| 4.15 | Batch Training | 9 |
| 4.16 | Understanding the Sigmoid | 9 |

4.1 Generative Classifiers

If we are [distinguishing](#) cat from dog images using a [Generative Classifier](#), we [build](#) a [model](#) of what is in a [cat image](#).

- Knows about whiskers, ears, eyes.
- [Assigns a probability](#) to any image to determine how cat-like is that image?

Similarly, [build](#) a [model](#) of what is in a [dog image](#). Now given a new image, run both models and see which one [fits better](#).

4.2 Discriminative Classifiers

If we are [distinguishing](#) cat from dog images using a [Discriminative Classifier](#).

- Just try to [distinguish](#) dogs from cats.

- Oh look, dogs have collars.
- Ignore everything else.

4.3 Generative vs Discriminative Classifiers

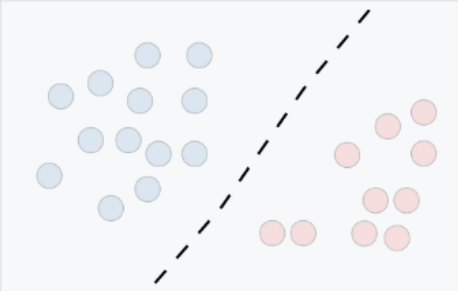
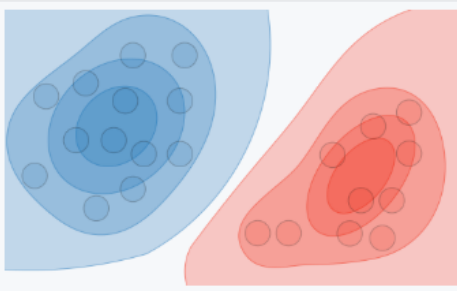
| | Discriminative model | Generative model |
|----------------|---|--|
| Goal | Directly estimate $P(y x)$ | Estimate $P(x y)$ to then deduce $P(y x)$ |
| What's learned | Decision boundary | Probability distributions of the data |
| Illustration |  |  |
| Examples | Regressions, SVMs | GDA, Naive Bayes |

Figure 4.1: Differences between Generative and Discriminative classifiers.

Generative Classifiers ([Naïve Bayes](#)) –

- Assume some functional form for [conditional independence](#).
- Estimate parameters of $P(D|h)$, $P(h)$ directly from training data.
- Use Bayes' rule to calculate $P(h|D)$.

Why not [learn](#) $P(h|D)$ or the decision boundary [directly](#)? Discriminative Classifiers ([Logistic Regression](#)) –

- Assume some functional form for $P(h|D)$ or for the decision boundary.
- [Estimate parameters](#) of $P(h|D)$ [directly](#) from training data.

4.4 Learning a Logistic Regression Classifier

Given n [input-output](#) pairs –

1. A feature representation of the [input](#). For each [input observation](#) x_i , a vector of [features](#) $[x_1, x_2, \dots, x_d]$.
2. A [classification function](#) that computes y , the estimated class, via $P(y|x)$, using the [sigmoid](#) or [softmax](#) functions.

3. An objective function for learning, like [cross-entropy loss](#).
4. An algorithm for optimizing the objective function, like [stochastic gradient ascent/descent](#).

4.5 Logistic Regression

Logistic Regression [assumes](#) the following function form for $P(y|x)$:

$$\begin{aligned}
 P(y = 1|x) &= \frac{1}{1 + e^{-(\sum_i w_i x_i + b)}} \\
 P(y = 1|x) &= \frac{1}{1 + e^{-(\sum_i w_i x_i + b)}} \\
 &= \frac{e^{(\sum_i w_i x_i + b)}}{e^{(\sum_i w_i x_i + b)} + 1} \\
 P(y = 0|x) &= 1 - \frac{1}{1 + e^{(\sum_i w_i x_i + b)}} \\
 &= \frac{1}{e^{(\sum_i w_i x_i + b)} + 1} \\
 \frac{P(y = 1|x)}{P(y = 0|x)} &= e^{(\sum_i w_i x_i + b)} > 1 \\
 &\Rightarrow \sum_i w_i x_i + b > 0
 \end{aligned}$$

Logistic Regression is a [linear](#) classifier. Turning a probability into a classifier using the [logistic function](#):

$$y_{LR} \begin{cases} 1 & \text{if } P(y = 1|x) \geq 0.5 \\ 0 & \text{otherwise} \end{cases} \begin{cases} \leftarrow w_i x_i + b \geq 0 \\ \leftarrow w_i x_i + b < 0 \end{cases}$$

4.6 LR Example

Suppose we are doing [binary sentiment classification](#) on movie review text, and we would like to know whether to assign the sentiment class [positive = 1](#) or [negative = 0](#) to the following review:

It's hokey. There are virtually no surprises, and the writing is second-rate. So why was is so enjoyable? For one thing, the case is great. Another nice touch is the music. I was overcome with the urge to get off the couch and start dancing. It sucked me in, and it'll do the same to you.

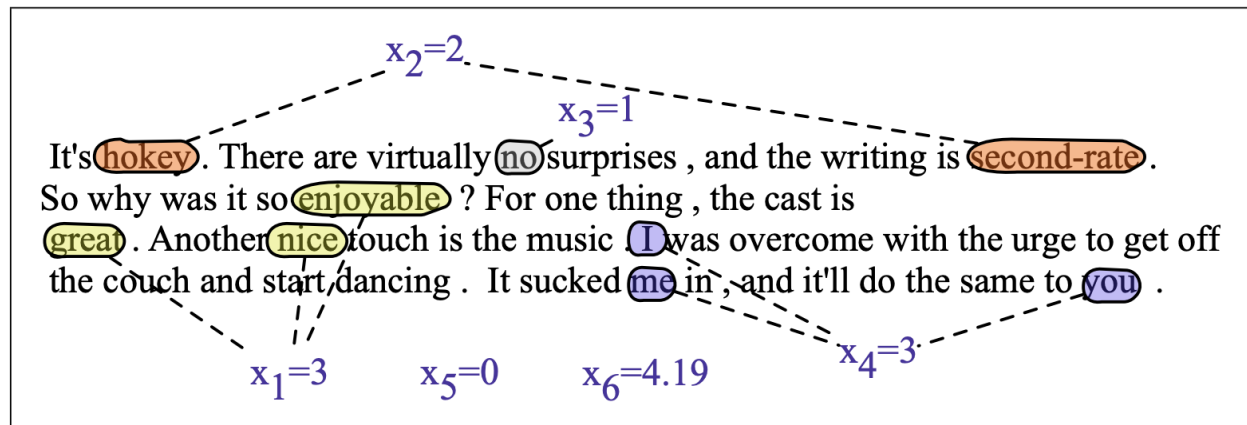


Figure 4.2: LR Example.

| | | |
|-------|---|------------------|
| x_1 | count(positive lexicon words \in doc) | 3 |
| x_2 | count(negative lexicon words \in doc) | 2 |
| x_3 | $\begin{cases} 1 & \text{if "no" } \in \text{ doc} \\ 0 & \text{otherwise} \end{cases}$ | 1 |
| x_4 | count(1st and 2nd pronouns \in doc) | 3 |
| x_5 | $\begin{cases} 1 & \text{if "!" } \in \text{ doc} \\ 0 & \text{otherwise} \end{cases}$ | 0 |
| x_6 | $\ln(\text{word count of doc})$ | $\ln(66) = 4.19$ |

Figure 4.3: Feature vector for the LR Example.

4.7 Sentiment Classification

Let's assume for the moment that we've already **learned** a **real-valued weight** for each of these features, and that the **6 weights** corresponding to the **6 features** are **[2.5, -5.0, -1.2, 0.5, 2.0, 0.7]**, while **$b = 0.1$** .

$$\begin{aligned}
P(+ve|x) &= P(y = 1|x) \\
&= \frac{1}{1 + e^{(\sum_i w_i x_i + b)}} \\
&= \frac{1}{1 + e^{-(2.5(3) + (-5)(2) + (-1.2)(1) + 0.5(3) + 2.0(0) + 0.7(4.19) + 0.1)}} \\
&= 0.30 \\
P(-ve|x) &= P(y = 0|x) \\
&= 1 - P(y = 1|x) \\
&= 1 - 0.70 \\
&= 0.30
\end{aligned}$$

Since $P(+ve|x) > P(-ve|x)$, the output sentiment class is **positive**.

4.8 Training Logistic Regression

We'll focus on **binary classification**. We **parameterize** (w_i, b) as θ :

$$\begin{aligned}
P(y_i = 0|x_i, \theta) &= \frac{1}{e^{\sum_i w_i x_i + b} + 1} \\
P(y_i = 1|x_i, \theta) &= \frac{e^{\sum_i w_i x_i + b}}{e^{\sum_i w_i x_i + b} + 1} \\
P(y_i|x_i, \theta) &= \frac{e^{y_i \sum_i w_i x_i + b}}{e^{\sum_i w_i x_i + b} + 1}
\end{aligned}$$

How do we **learn parameters** θ ?

4.9 Cross-Entropy Loss

- We want to know **how far** is the **classifier output** \hat{y} from the **true output** y . Let's call this difference $L(\hat{y}, y)$.
- Since there are only **2 discrete outcomes** (0 or 1), we can express the probability $P(y|x)$ from our classifiers as:

$$P(y|x) = \hat{y}^y \cdot (1 - \hat{y})^{1-y}$$

- Goal: **maximize the probability** of the correct label $P(y|x)$.
- Maximize:

$$\begin{aligned}
P(y|x) &= \hat{y}^y \cdot (1 - \hat{y})^{1-y} \\
\log(P(y|x)) &= \log(\hat{y}^y \cdot (1 - \hat{y})^{1-y}) \\
&= y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})
\end{aligned}$$

- We want to **minimize** the **cross-entropy loss**:

$$\begin{aligned}
 \text{Minimize : } L_{CE}(\hat{y}, y) &= -\log P(y|x) \\
 &= -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})] \\
 \min_{\theta} L_{CE}(\hat{y}, y) &= -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})] \\
 &= -\left[y \log\left(\frac{e^{\sum_i w_i x_i + b}}{1 + e^{\sum_i w_i x_i + b}}\right) + (1 - y) \log\left(1 - \frac{e^{\sum_i w_i x_i + b}}{1 + e^{\sum_i w_i x_i + b}}\right) \right] \\
 &= -\left[y \left(\sum_i w_i x_i + b - \log(1 + e^{\sum_i w_i x_i + b}) \right) + (1 - y) \left(-\log(1 + e^{\sum_i w_i x_i + b}) \right) \right] \\
 &= -\left[y \left(\sum_i w_i x_i + b \right) - \log(1 + e^{\sum_i w_i x_i + b}) \right] \\
 &= \log(1 + e^{\sum_i w_i x_i + b}) + y \left(\sum_i w_i x_i + b \right)
 \end{aligned}$$

4.10 Minimizing Cross-Entropy Loss

$$\min_{\theta} L_{CE}(\hat{y}, y)$$

- Minimizing loss function $L_{CE}(\hat{y}, y)$ is a **convex optimization** problem.
- **Convex** function have a **global minimum**.
- **Concave** function have a **global maxima**.

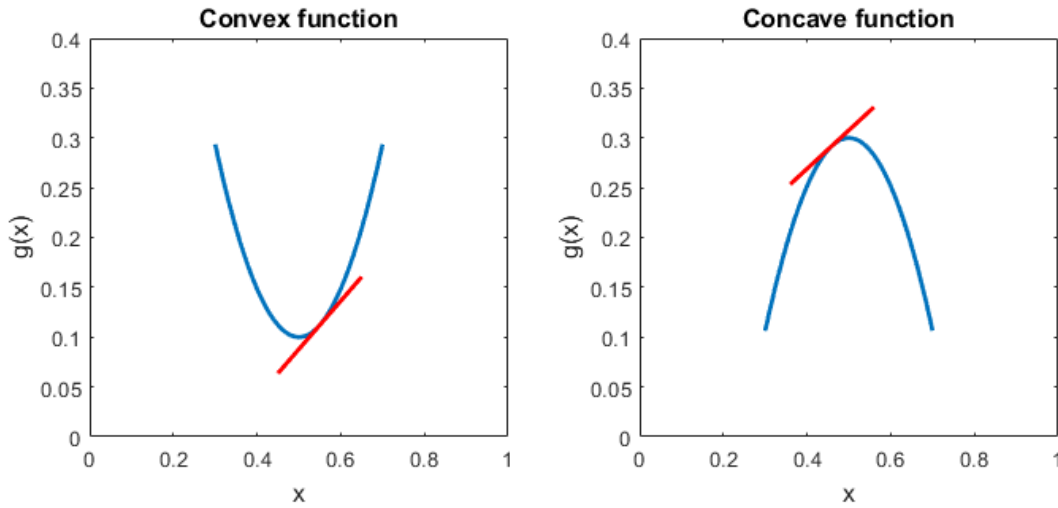


Figure 4.4: An example of a convex and concave function.

4.11 Optimizing a Convex/Concave Function

- Maximum of a **concave** function is **equivalent** to the **minimum** of a **convex** function.
- **Gradient Ascent** is used for finding the **maximum** of a **concave** function.
- **Gradient Descent** is used for finding the **minimum** of a **convex** function.

4.12 Gradients

- The **gradient** of a **function** is a **vector pointing** in the **direction** of the **greatest increase** in a function.

Gradient Ascent: Find the gradient of the function at the current point and **move** in the **same direction**.

Gradient Descent: Find the gradient of the function at the current point and **move** in the **opposite direction**.

4.13 Gradient Descent for Logistic Regression

- Let us represent $\hat{y} = f(x, \theta)$
- Gradient:

$$\nabla_{\theta} L(f(x, \theta), y) = \left[\frac{\partial L(f(x, \theta), y)}{\partial b}, \frac{\partial L(f(x, \theta), y)}{\partial w_1}, \frac{\partial L(f(x, \theta), y)}{\partial w_2}, \dots, \frac{\partial L(f(x, \theta), y)}{\partial w_d} \right] \quad (4.1)$$

- Update Rule:

$$\begin{aligned} \Delta \theta &= \eta \cdot \nabla_{\theta} L(f(x, \theta), y) \\ \theta_{t+1} &= \theta_t - \eta \cdot \frac{\partial}{\partial (w, b)} L(f(x, \theta), y) \end{aligned} \quad (4.2)$$

Gradient descent algorithm will **iterate** until $\Delta \theta < \epsilon$.

$$\begin{aligned} L_{CE}(f(x, \theta), y) &= \log(1 + e^{\sum_i w_i x_i + b}) - y \left(\sum_i w_i x_i + b \right) \\ \theta_{t+1} &= \theta_t - \eta \cdot \frac{\partial}{\partial (w, b)} L(f(x, \theta), y) \\ &= \theta_t - \eta \cdot x_i \left[\frac{e^{\sum_i w_i x_i + b}}{1 + e^{\sum_i w_i x_i + b}} - y \right] \\ &= \theta_t - \eta \cdot x_i \left[\hat{P}(y = 1 | x, \theta_t) - y \right] \end{aligned}$$

4.14 Learning Rate

- η is a **hyperparameter**.
- **Large η** \Rightarrow Fast convergence but larger residual error. Also, possible oscillations.
- **Small η** \Rightarrow Slow convergence but small residual error.

4.15 Batch Training

- Stochastic gradient descent is called **stochastic** because it chooses a **single random example** at a time, moving the weights to improve performance on that single example.
- This results in very choppy movements, so it's **common** to **compute** the **gradient over batches** of training instances rather than a single instance.
- **Training data:** $\{x_i, y_i\}_{i=1 \dots n}$ where $x_i = (x_{i1}, x_{i2}, \dots, x_{id})$, n is the total **instances** in a **batch** and d is the **dimension** of an instance.

$$\theta_{t+1} = \theta_t - \frac{\eta}{n} \times \sum_{i=1}^n x_i j \left[\frac{1}{1 + e^{-\theta^T x}} - y_i \right] \quad (4.3)$$

4.16 Understanding the Sigmoid

- Large weights lead to **overfitting**.
- **Penalizing** larger weights can **reduce** overfitting.