

Chapter 4

Logistic Regression

Contents

4	Logistic Regression	1
4.1	Generative Classifiers	2
4.2	Discriminative Classifiers	2
4.3	Generative vs Discriminative Classifiers	3
4.4	Learning a Logistic Regression Classifier	3
4.5	Logistic Regression	3
4.6	LR Example	4
4.7	Sentiment Classification	4
4.8	Training Logistic Regression	4
4.9	Cross-Entropy Loss	4
4.10	Minimizing Cross-Entropy Loss	5
4.11	Optimizing a Convex/Concave Function	5
4.12	Gradients	5
4.13	Gradient Descent for Logistic Regression	5
4.14	Learning Rate	6

4.1 Generative Classifiers

If we are [distinguishing](#) cat from dog images using a [Generative Classifier](#), we [build](#) a [model](#) of what is in a [cat image](#).

- Knows about whiskers, ears, eyes.
- [Assigns a probability to any image to determine how cat-like is that image?](#)

Similarly, [build](#) a [model](#) of what is in a [dog image](#). Now given a new image, run both models and see which one [fits better](#).

4.2 Discriminative Classifiers

If we are [distinguishing](#) cat from dog images using a [Discriminative Classifier](#).

- Just try to [distinguish](#) dogs from cats.
 - Oh look, dogs have collars.
 - Ignore everything else.

4.3 Generative vs Discriminative Classifiers

Generative Classifiers (Naïve Bayes) –

- Assume some functional form for conditional independence.
- Estimate parameters of $P(D|h)$, $P(h)$ directly from training data.
- Use Bayes rule to calculate $P(h|D)$.

Why not learn

4.4 Learning a Logistic Regression Classifier

Given n input-output pairs –

1. A feature representation of the input. For each input observation x_i , a vector of features $[x_1, x_2, \dots, x_d]$.
2. A classification function that computes y , the estimated class, via $P(y|x)$, using the sigmoid of softmax functions.
3. An objective function for learning, like cross-entropy loss.
4. An algorithm for optimizing the objective function, like stochastic gradient ascent/descent.

4.5 Logistic Regression

Logistic Regression assumes the following function form for $P(y|x)$:

$$\begin{aligned}
 P(y = 1|x) &= \frac{1}{1 + e^{-(\sum_i w_i x_i + b)}} \\
 P(y = 1|x) &= \frac{1}{1 + e^{-(\sum_i w_i x_i + b)}} \\
 &= \frac{e^{(\sum_i w_i x_i + b)}}{e^{(\sum_i w_i x_i + b)} + 1} \\
 P(y = 0|x) &= 1 - \frac{1}{1 + e^{(\sum_i w_i x_i + b)}} \\
 &= \frac{1}{e^{(\sum_i w_i x_i + b)} + 1} \\
 \frac{P(y = 1|x)}{P(y = 0|x)} &= e^{(\sum_i w_i x_i + b)} > 1 \\
 &\Rightarrow \sum_i w_i x_i + b > 0
 \end{aligned}$$

Logistic Regression is a **linear** classifier. Turning a probability into a classifier using the **logistic function**:

$$y_{LR} \begin{cases} 1 & \text{if } P(y = 1|x) \geq 0.5 \\ 0 & \text{otherwise} \end{cases} \begin{matrix} \leftarrow w_i x_i + b \geq 0 \\ \leftarrow w_i x_i + b < 0 \end{matrix}$$

4.6 LR Example

4.7 Sentiment Classification

Let's assume for the moment that we've already **learned** a **real-valued weight** for each of these features, and that the **6 weights** corresponding to the **6 features** are $[2.5, -5.0, -1.2, 0.5, 2.0, 0.7]$, while $b = 0.1$.

4.8 Training Logistic Regression

We'll focus on **binary classification**. We **parameterize** w_i, b as θ :

$$P(y_i = 0|x_i, \theta) = \frac{1}{e^{\sum_i w_i x_i + b} + 1}, P(y_i = 1|x_i, \theta) = \frac{e^{\sum_i w_i x_i + b}}{e^{\sum_i w_i x_i + b} + 1}, P(y_i|x_i, \theta) = \frac{e^{y_i \sum_i w_i x_i + b}}{e^{\sum_i w_i x_i + b} + 1}$$

How do we **learn parameters** θ ?

4.9 Cross-Entropy Loss

- We want to know **how far** is the **classifier output** \hat{y} from the **true output** y . Let's call this difference $L(\hat{y}, y)$.
- Since there are only **2 discrete outcomes** (0 or 1), we can express the probability $P(y|x)$ from our classifiers as:

$$P(y|x) = \hat{y}^y \cdot (1 - \hat{y})^{1-y}$$

- Goal: **maximize the probability** of the correct label $P(y|x)$.
- Maximize:

$$\begin{aligned} P(y|x) &= \hat{y}^y \cdot (1 - \hat{y})^{1-y} \\ \log(P(y|x)) &= \log(\hat{y}^y \cdot (1 - \hat{y})^{1-y}) \\ &= y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}) \end{aligned}$$

- We want to **minimize** the **cross-entropy loss**:

4.10 Minimizing Cross-Entropy Loss

$$\min_{\theta} L_{CE}(\hat{y}, y)$$

- Minimizing loss function $L_{CE}(\hat{y}, y)$ is a **convex optimization problem**.

4.11 Optimizing a Convex/Concave Function

- Maximum of a **concave** function is **equivalent** to the **minimum** of a **convex** function.
- Gradient Ascent** is used for finding the **maximum** of a **concave** function.
- Gradient Descent** is used for finding the **minimum** of a **convex** function.

4.12 Gradients

- The **gradient** of a **function** is a **vector pointing** in the **direction** of the **greatest increase** in a function.

Gradient Ascent

Gradient Descent

4.13 Gradient Descent for Logistic Regression

- Let us represent $\hat{y} = f(x, \theta)$
- Gradient:

$$\nabla_{\theta} L(f(x, \theta), y) = \left[\frac{\partial L(f(x, \theta), y)}{\partial b}, \frac{\partial L(f(x, \theta), y)}{\partial w_1}, \frac{\partial L(f(x, \theta), y)}{\partial w_2}, \dots, \frac{\partial L(f(x, \theta), y)}{\partial w_d} \right] \quad (4.1)$$

- Update Rule:

$$\begin{aligned} \Delta \theta &= \eta \cdot \nabla_{\theta} L(f(x, \theta), y) \\ \theta_{t+1} &= \theta_t - \eta \cdot \frac{\partial}{\partial (w, b)} L(f(x, \theta), y) \end{aligned} \quad (4.2)$$

Gradient descent algorithm will **iterate** until $\Delta \theta < \epsilon$.

$$\begin{aligned} L_{CE}(f(x, \theta), y) &= \log(1 + e^{\sum_i w_i x_i + b}) - y \left(\sum_i w_i x_i + b \right) \\ \theta_{t+1} &= \theta_t - \eta \cdot \frac{\partial}{\partial(w, b)} L(f(x, \theta), y) \\ &= \theta_t - \eta \cdot x_i \left[\frac{e^{\sum_i w_i x_i + b}}{1 + e^{\sum_i w_i x_i + b}} - y \right] \\ &= \theta_t - \eta \cdot x_i \left[\hat{P}(y = 1 | x, \theta_t) - y \right] \end{aligned}$$

4.14 Learning Rate

- η is a [hyperparameter](#).
- [Large](#) $\eta \Rightarrow$ Fast convergence but larger residual error. Also, possible oscillations.
- [Small](#) $\eta \Rightarrow$ Slow convergence but small residual error.