

Chapter 8

Artificial Neural Network

Contents

8 Artificial Neural Network	1
8.1 ANN Properties	2
8.2 Perceptron	2
8.3 Multi-layer Perceptron	3
8.4 Activation Functions	3
8.5 Feed-forward MLP	4
8.6 Forward Pass	5
8.7 Gradient Descent	5
8.8 Chain Rule	5
8.9 Backward Pass	5
8.10 Backpropagation using Sigmoid	6
8.11 Backpropagation Example	6
8.12 Backpropagation with Batch Gradient Descent	7
8.13 Training	7
8.14 Dropout	7
8.15 Exploding Gradient Problem	8
8.16 Dying ReLU Problem	8

8.1 ANN Properties

- ANN is much like a [black-box](#).
- Many [neuron-like](#) threshold [switching](#) units.
- Many [weighted interconnections](#) amount units.
- Highly [parallel](#) and [distributed](#) process.

8.2 Perceptron

Linear function: $f : X \rightarrow Y$

$$\begin{aligned} f(\mathbf{x}) &= w_0 + w_1x_1 + w_2x_2 + \cdots + w_dx_d \\ &= w_0 + \sum_{j=1}^d w_jx_j \end{aligned}$$

w_0, w_1, \dots, w_d are weights

8.3 Multi-layer Perceptron

- Multi-layer perceptrons (MLPs) were designed to overcome the computational limitation of a single threshold element (perceptron).
- The idea is to stack several layers of threshold elements, each layer using the output of the previous layer as input.
- A feed-forward MLP network defines a mapping:

$$y = f(x; \theta)$$

- Functions are composed in a chain:

$$f(x) = f_3(f_2(f_1(x)))$$

- **Input layer:** The process starts with the input layer, which receives the input data. Each neuron in this layer represents a feature of input data.
- **Weights and Biases:** Connections between neurons have associated weights, which are learned during the training process and is crucial to capture patterns in the data.
- **Hidden Layers:** The neurons in these layers perform computations on the inputs. The output of each neuron is calculated by applying a weighted sum of its inputs (from the previous layer).
- **Activation Functions:** The activation function is crucial as it introduces non-linearity into the model, allowing it to learn more complex functions. ...

8.4 Activation Functions

- **Non-linearity:** This is the most fundamental property; activation functions introduce nonlinearity into the network. This is important because real-world relationships and patterns are rarely linear.
- ...
- **Monotonicity:** A monotonic activation function either strictly increases or strictly decreases as input values change. This property ensures that as inputs change, the neuron's output moves in a consistent direction.
- **Continuity:** A continuous activation function produces smooth and continuous changes in output as inputs change slightly. This property helps in smooth gradient computations for updating weights during the learning process.

- **Differentiability**: Differentiability is essential for **gradient-based optimization** algorithms like backpropagation. **Activation functions** that are **differentiable** across their domain allow **gradients** to be computed for **weight updates** during training.
- **Sparsity**: Some activation functions promote sparsity by having their **outputs** be **zero** for a large portion of input space. This can be beneficial in **reducing** the **complexity** of neural networks.

Rectified Linear Units (ReLU):

$$\begin{aligned} g(x) &= \max\{0, x\} \\ g'(x) &= \begin{cases} 1, & x \geq 0 \\ 0, & x < 0 \end{cases} \end{aligned} \quad (8.1)$$

Sigmoid:

$$\begin{aligned} g(x) &= \frac{1}{1 + e^{-x}} \\ g'(x) &= g(x)(1 - g(x)) \end{aligned} \quad (8.2)$$

Hyperbolic tangent tanh

$$\begin{aligned} g(x) &= \frac{e^x - e^{-x}}{e^x + e^{-x}} \\ g'(x) &= 1 - g(x)^2 \end{aligned} \quad (8.3)$$

Table 8.1: Activation Functions

Activation Functions	Sigmoid	tanh	ReLU
Range	(0, 1)	(-1, 1)	[0, ∞)
Vanishing Gradient Problem	Yes	Yes	No
Nature	Non-Linear	Non-Linear	Linear
Zero Centered Activation Function	No	Yes	No
Symmetric Function	No	Yes	No

8.5 Feed-forward MLP

- The **weights** of the **neural network connections** are **repeatedly adjusted** to **minimize** the **difference** between the **actual output** and **desired output**.
- Aims to **minimize** the **loss function** by **adjusting** the **network's weights and biases**. The **loss function gradients** determine the level of adjustment with respect to parameters like activation function, weights, bias, etc.
- The **forward step**, given the input, **computes** the **output layer-by-layer**, starting with the **input layer**.
- The **backward step** **calculates** the **error** in the output and **propagates** it **backwards**; then **update** the **weights layer-by-layer**, starting from the output layer.

8.6 Forward Pass

- In this example, we will be using a [three-layer neural network](#) with the layers being the input, hidden, and output layers.
- The [activation](#) ...

8.7 Gradient Descent

- Goal of NN: Given n training samples (x_i, y_i) , find w to [minimize](#):

$$E[w] = \frac{1}{2n} \sum_{i=1}^n (y_i - o_i)^2 \quad (8.4)$$

- ...

8.8 Chain Rule

- $\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x}$
- $\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y_1} \frac{\partial y_1}{\partial x} + \frac{\partial z}{\partial y_2} \frac{\partial y_2}{\partial x}$

8.9 Backward Pass

- Function 1 ([error](#)):

$$\begin{aligned} \frac{\partial E_h}{\partial o_h} &= -(y_h - o_h) \\ &= o_h - y_h \end{aligned}$$

- Function 2 ([differentiable activation function](#)):

$$\begin{aligned} \frac{\partial o_h}{\partial net_h} &= \frac{e^{-net_h}}{(1 + e^{-net_h})^2} \\ &= \frac{1}{1 + e^{-net_h}} \times \left(1 - \frac{1}{1 + e^{-net_h}}\right) \\ &= o_h(1 - o_h) \end{aligned}$$

- Function 3 ([linear](#) gate):

$$\frac{\partial net_h}{\partial w_{jh}} = x_j$$

Output neuron backpropagation error:

$$\begin{aligned}\delta_1^y &= \frac{\partial E_1}{\partial o_1^y} \frac{\partial o_1^y}{\partial net_1} \\ &= (o_1^y - y_1) o_1^y (1 - o_1^y)\end{aligned}\tag{8.5}$$

Hidden neuron backpropagation error:

$$\delta_1^{h_2} = \delta_1^{h_2} (1 - o_1^{h_2}) \sum_{m=1}^3 \delta_m^y w_{h\frac{1}{2}y_m}\tag{8.6}$$

Weight updates:

$$\begin{aligned}w_{h\frac{1}{2}y_1} &= w_{h\frac{1}{2}y_1} - \eta \delta_1^y o_1^{h_2} \\ w_{h\frac{1}{1}h\frac{1}{2}} &= w_{h\frac{1}{1}h\frac{1}{2}} - \eta \delta_1^{h_2} o_1^{h_1}\end{aligned}$$

8.10 Backpropagation using Sigmoid

Repeat for each training example until end of training epoch:

1. In the forward pass, compute the **outputs** o of each neuron
 - (a) The outputs for the **input layer neurons** stay unchanged.
 - (b) The outputs for the **hidden** and **output layer neurons** are computed using the **sigmoid** activation function.
2. In the **backward pass**, **propagate** the **errors** δ from **output layer**
 - (a) For each **output** neuron k : $\delta_k = (o_k - y_k) o_k (1 - o_k)$
 - (b) For each **hidden** neuron: h : $\delta_h = o_h (1 - o_h) \sum_{m=1}^{|h'|} \delta_m w_{hm}$, where h' is the **subsequent layer**.
3. Update every **weights** $w_{ij} \leftarrow w_{ij} - \eta \delta_j o_i$ where w_{ij} denotes the **weight** between **nodes** i and j .

8.11 Backpropagation Example

- x are the **inputs**, h are the **neurons** with **sigmoid activation**, y are the **outputs**.
- The **initial weights** of all the layers are shown in the figure.
- Our objective is to learn ... $\eta = 1$.

8.12 Backpropagation with Batch Gradient Descent

- The [backpropagation](#) algorithm we covered so far, as well as the example we practiced, both [used stochastic gradient descent](#) for weight updates i.e., the [weights](#) are [updated](#) for [every training instances](#).
- In [backpropagation](#) using bgd, the [weights](#) are [updated](#) ...

8.13 Training

[Backpropagation](#) is a gradient estimation method used to train NNs –

- [No guarantee](#) of [convergence](#) since NNs form [non-convex functions](#) with [multiple local minima](#).
- [Many epochs](#) (tens of thousands) may be needed for adequate training.
- [Large data sets](#) may require many hours of CPU.
- [Termination criteria](#): Number of epochs, threshold on training set error, early stopping, increased error on validation set.
-
- [Underfitting](#) –
 - Using [too few](#) hidden neurons in the NN.
 - [Inadequate](#) or less [data](#) to train the NN on.
- [Overfitting](#) –
 - Training the NN over [too many epochs](#).
 - Using [too many hidden layers](#) in the NN.
 - Using [too many neurons](#) in a hidden layer in the NN.

8.14 Dropout

- [Large weights](#) in NN are a sign or a more [complex network](#) that has overfit the training data.
- Probabilistically [dropping out nodes](#) in the network is a simple and effective regularization method.

8.15 Exploding Gradient Problem

8.16 Dying ReLU Problem

- A **dying RELU** always **outputs** the same value i.e., **0** on any input value.
- This condition is known as the **dead state** of the ReLU **neurons**.
- In this state, it is **difficult to recover** because the **gradient** of **0** is **0**.
- This becomes a problem when most of the **input ranges** are **negative**, or the **derivative** of the **ReLU function** is **0**.
- Can be caused by a **high learning rate** or a **large negative bias**.
- Using **Leaky ReLU** instead can resolve the dying ReLU problem.

$$\begin{aligned} g(x) &= \max\{0.01x, x\} \\ g'(x) &= \begin{cases} 1, & x \geq 0 \\ 0.01, & x < 0 \end{cases} \end{aligned} \tag{8.7}$$