

Chapter 10

Recurrent Neural Networks

Contents

10 Recurrent Neural Networks	1
10.1 Motivation	2
10.2 Neurons with Recurrent	2
10.3 Recurrent Neural Network	3
10.4 RNN Implementation	3
10.5 Unrolling the RNN	3
10.6 Sequence Modeling	3
10.7 Backpropagation Through Time	4
10.8 LSTM Cell	4
10.9 LSTM Memory	4
10.10 Forget Gate	4
10.11 Input Gate	4
10.12 Updating Cell State	5
10.13 Output Gate	5
10.14 Training LSTM	5
10.15 LSTM BPTT	5
10.16 LSTM Example	6
10.17 Gated Recurrent Unit	6
10.18 Bidirectional RNN	6
10.19 Seq2Seq	6
10.20 Limitations of RNN	6
10.21 Desired Capabilities of RNN	6

10.1 Motivation

- When the model generates “people”, we need a way to [tell the model](#) that “several” has [already been generated](#) and similarly for the other words.

10.2 Neurons with Recurrent

$$y^t = f(x^t, h^{t-1}) \tag{10.1}$$

10.3 Recurrent Neural Network

- Apply a [recurrence relation](#) at every [time step](#) t to [process](#) a [sequence](#).
- The [feedback connection](#) allows [information](#) to [persist](#).
- [RNNs](#) have [hidden state](#) h^t

10.4 RNN Implementation

- ...
- [Loop](#) through all [individual words](#) in the sentence.
- Inside the loop, each [word](#) is [fed](#) into [RNN](#) model along with [previous hidden state](#).
- This [generates](#) ...
- [Prediction](#) for the [final word](#) is the [RNN's output](#) after all the [prior work](#) have been [fed](#) in [through](#) the [model](#).

10.5 Unrolling the RNN

$$\begin{aligned} h^t &= g(w_{xh}x^t + w_{hh}h^{t-1}) \\ y^t &= g(w_{hy}h^t) \end{aligned} \tag{10.2}$$

$$L = \sum_{t=0}^k L^t \tag{10.3}$$

10.6 Sequence Modeling

- [Sequence models](#) have been [motivated](#) by the [analysis](#) of [sequential data](#) such as text sentences, time-series and other discrete sequences data.
- These [models](#) are designed to [handle sequential information](#) in the same way that CNN are adapted to handle spatial data.
- The [key point](#) for [sequence models](#) is that the [data](#) we are processing are [not independently and identically distributed](#) samples and the data carry some [dependency](#) due to their [sequential ordering](#).
- [Applications](#) of sequence models ...
- Sequence modeling [design criteria](#):

- Handle [variable-length sequences](#).
 - Track [long-term dependencies](#).
 - Maintain [information](#) about the [order](#).
 - [Share parameters](#) across the sequence.
- RNNs satisfy all of the design criteria for sequence modeling.

10.7 Backpropagation Through Time

10.8 LSTM Cell

10.9 LSTM Memory

- The [key](#) of an LSTM is the [cell state](#), the [horizontal line](#) running through the top of the diagram.
- The [cell state](#) is like a [conveyor belt](#). It's very [easy](#) for [information](#) to just [flow](#) along it [unchanged](#).

10.10 Forget Gate

- It looks at h^{t-1} and x^t , and [outputs](#) a number between [0](#) and [1](#) for each number in the [cell state](#) c^t .
- [1](#) represents “completely keep this” while a [0](#) represents “completely get rid of this”.

$$f^t = \sigma(w_{xf}x^t + w_{hf}h^{t-1} + b_f) \quad (10.4)$$

10.11 Input Gate

- First, a [sigmoid layer](#) called the “[input gate layer](#)” decides [which values to update](#).
- Second, a [tanh layer](#) creates a vector of [new candidate values](#) c^t , that could be [added](#) to the [state](#).

$$\begin{aligned} i^t &= \sigma(w_{xi}x^t + w_{hi}h^{t-1} + b_i) \\ c^t &= \tanh(w_{xc}x^t + w_{hc}h^{t-1} + b_c) \end{aligned} \quad (10.5)$$

10.12 Updating Cell State

- Next, **update** the **old cell state** c^{t-1} into the new cell state c^t .
- We use the **Hadamard product** \circ (**element-wise** product).

$$c^t = f^t \circ c^{t-1} + i^t \circ c^t \quad (10.6)$$

10.13 Output Gate

- A **sigmoid layer decides** what **parts** of the **cell state** are going to be **output**.
- The **cell state** is then put through **tanh** (to make the values $\in [-1, 1]$) and this is **element-wise multiplied** by the output of the sigmoid gate.

$$o^t \quad (10.7)$$

10.14 Training LSTM

- An LSTM network is **trained** with **BPTT**.
- In a **vanilla RNN**, if $h^t = (w_{xh}x^t + w_{hh}h^{t-1})$, then its **derivative** is –

$$\frac{\partial h^t}{\partial h^{t-1}} = w_{hh}\sigma(\cdot)(1 - \sigma(\cdot))$$

- However, for **LSTM**, we have the **cell state** $c^t = f^t \circ c^{t-1} + i^t \circ c^t$ and its **derivative** is –

$$\frac{\partial c^t}{\partial c^{t-1}} = f^t = \sigma.$$

- This helps **LSTM preserve** a **constant error** when it is back-propagated.
- The **cells learn** when to allow data to **enter, leave or delete** through the iterative process of **back-propagating errors** and **adjusting weights**.

10.15 LSTM BPTT

$$o^t = \sigma(w_{xo}x^t + w_{ho})$$

10.16 LSTM Example

- Cell state and hidden state represents long-term and short-term memory.
- Forget gate will decide how much of the long-term memory to pass on.

10.17 Gated Recurrent Unit

- Similar performance as LSTM with less computation.
- GRUs have fewer parameters than LSTM, as they lack an output gate.

10.18 Bidirectional RNN

- Output at time t may not only depend on the previous elements in the sequence, but also future elements.
- They are just two RNNs stacked on top of each other. The output is then computed based on the hidden state of both RNNs.

10.19 Seq2Seq

- Consists of two RNNs:
 - The encoder reads the input sequence and outputs a vector.
 - The decoder reads the vector and produces the output sequence.
- Primarily used in NLP applications.

10.20 Limitations of RNN

- Encoding bottleneck
- Slow, no parallelization
- Not long memory

10.21 Desired Capabilities of RNN

- Continuous stream
- Parallelization
- Long memory